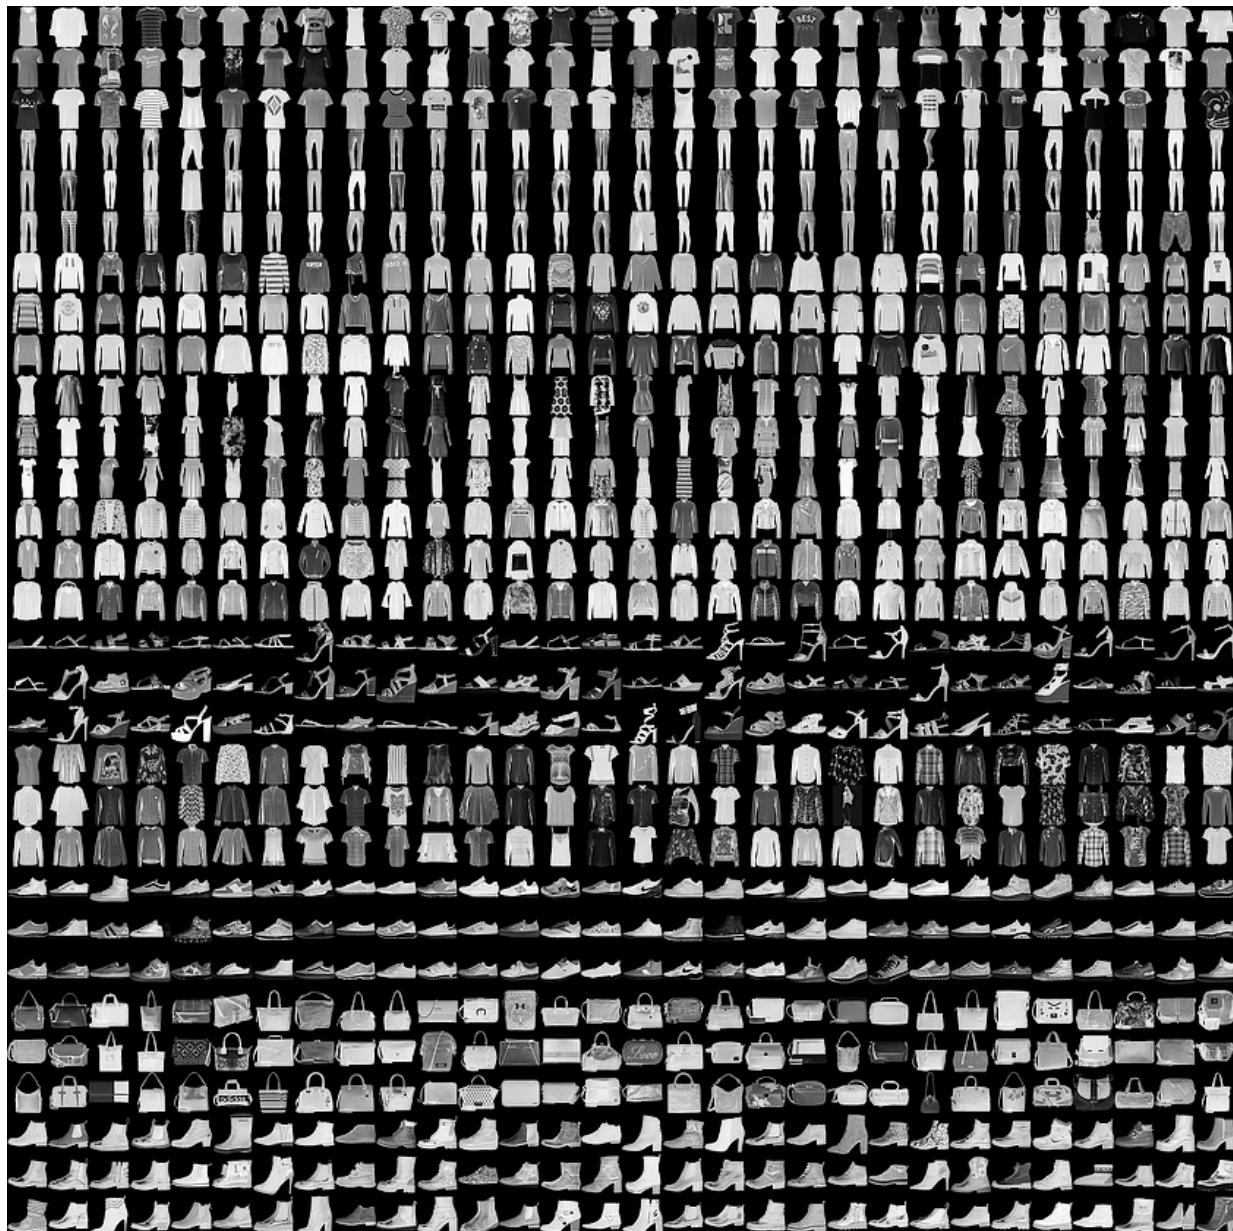


[zalandoresearch / fashion-mnist](#)

A MNIST-like fashion product database. Benchmark  <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/>

#mnist #deep-learning #benchmark #machine-learning #dataset #computer-vision #fashion #fashion-mnist #gan #zalando #convolutional-neural-networks

 201 commits	 1 branch	 0 releases	 4 contributors	 MIT
Branch: master New pull request		Create new file Upload files Find file Clone or download		
 hanxiao Update README.ja.md Latest commit 629abc0 on Oct 5, 2018				
 benchmark	first commit			a year ago
 data	fixes #44 I removed duplicate samples from the data set			a year ago
 doc/img	update image			a year ago
 static	add highlight to column header			a year ago
 utils	simplified reader			a year ago
 visualization	first commit			a year ago
 .catwatch.yaml	add zappr and catwatch			a year ago
 .dockerignore	first commit			a year ago
 .gitignore	update readme			a year ago
 .zappr.yaml	add zappr and catwatch			a year ago
 CONTRIBUTING.md	first commit			a year ago
 Dockerfile	first commit			a year ago
 LICENSE	first commit			a year ago
 MAINTAINERS	Update MAINTAINERS			10 months ago
 README.ja.md	Update README.ja.md			3 months ago
 README.md	Update README.md			3 months ago
 README.zh-CN.md	Update README.zh-CN.md			3 months ago
 app.py	first commit			a year ago
 configs.py	first commit			a year ago
 requirements.txt	first commit			a year ago
 README.md				
<h2>Fashion-MNIST</h2> <hr/> <p> Star 5k chat on gitter README 中文 README 日本語 License MIT  Year in Review </p>				
<p>► Table of Contents</p> <p>Fashion-MNIST is a dataset of Zalando's article images—consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes. We intend Fashion-MNIST to serve as a direct drop-in replacement for the original MNIST dataset for benchmarking machine learning algorithms. It shares the same image size and structure of training and testing splits.</p> <p>Here's an example how the data looks (<i>each class takes three-rows</i>):</p>				





Why we made Fashion-MNIST

The original [MNIST dataset](#) contains a lot of handwritten digits. Members of the AI/ML/Data Science community love this dataset and use it as a benchmark to validate their algorithms. In fact, MNIST is often the first dataset researchers try. *"If it doesn't work on MNIST, it won't work at all"*, they said. *"Well, if it does work on MNIST, it may still fail on others."*

To Serious Machine Learning Researchers

Seriously, we are talking about replacing MNIST. Here are some good reasons:

- **MNIST is too easy.** Convolutional nets can achieve 99.7% on MNIST. Classic machine learning algorithms can also achieve 97% easily. Check out [our side-by-side benchmark for Fashion-MNIST vs. MNIST](#), and read "[Most pairs of MNIST digits can be distinguished pretty well by just one pixel.](#)"
- **MNIST is overused.** In [this April 2017 Twitter thread](#), Google Brain research scientist and deep learning expert Ian Goodfellow calls for people to move away from MNIST.
- **MNIST can not represent modern CV tasks,** as noted in [this April 2017 Twitter thread](#), deep learning expert/Keras author François Chollet.

Get the Data

[Many ML libraries](#) already include Fashion-MNIST data/API, give it a try!

You can use direct links to download the dataset. The data is stored in the **same** format as the original [MNIST data](#).

Name	Content	Examples	Size	Link	MD5 Checksum
train-images-idx3-ubyte.gz	training set images	60,000	26 MBytes	Download	8d4fb7e6c68d591d4c3dfef9ec88bf0d

Name	Content	Examples	Size	Link	MD5 Checksum
train-labels-idx1-ubyte.gz	training set labels	60,000	29 KBytes	Download	25c81989df183df01b3e8a0aad5dffbe
t10k-images-idx3-ubyte.gz	test set images	10,000	4.3 MBytes	Download	bef4ecab320f06d8554ea6380940ec79
t10k-labels-idx1-ubyte.gz	test set labels	10,000	5.1 KBytes	Download	bb300cfdad3c16e7a12a480ee83cd310

Alternatively, you can clone this GitHub repository; the dataset appears under `data/fashion`. This repo also contains some scripts for benchmark and visualization.

```
git clone git@github.com:zalandoresearch/fashion-mnist.git
```

Labels

Each training and test example is assigned to one of the following labels:

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

Usage

Loading data with Python (requires NumPy)

Use `utils/mnist_reader` in this repo:

```
import mnist_reader
X_train, y_train = mnist_reader.load_mnist('data/fashion', kind='train')
X_test, y_test = mnist_reader.load_mnist('data/fashion', kind='t10k')
```

Loading data with Tensorflow

Make sure you have [downloaded the data](#) and placed it in `data/fashion`. Otherwise, *Tensorflow will download and use the original MNIST*.

```
from tensorflow.examples.tutorials.mnist import input_data
data = input_data.read_data_sets('data/fashion')
```

```
data.train.next_batch(BATCH_SIZE)
```

Note, Tensorflow supports passing in a source url to the `read_data_sets`. You may use:

```
data = input_data.read_data_sets('data/fashion', source_url='http://fashion-mnist.s3-website.eu-central-1.amazonaws.co
```

Also, an official Tensorflow tutorial of using `tf.keras`, a high-level API to train Fashion-MNIST [can be found here](#).

Loading data with other machine learning libraries

To date, the following libraries have included `Fashion-MNIST` as a built-in dataset. Therefore, you don't need to download `Fashion-MNIST` by yourself. Just follow their API and you are ready to go.

- [Apache MXNet Gluon](#)
- [deeplearn.js](#)
- [Kaggle](#)
- [Pytorch](#)
- [Keras](#)
- [Edward](#)
- [Tensorflow](#)
- [Torch](#)
- [JuliaML](#)
- [Chainer](#)

You are welcome to make pull requests to other open-source machine learning packages, improving their support to `Fashion-MNIST` dataset.

Loading data with other languages

As one of the Machine Learning community's most popular datasets, MNIST has inspired people to implement loaders in many different languages. You can use these loaders with the `Fashion-MNIST` dataset as well. (Note: may require decompressing first.) To date, we haven't yet tested all of these loaders with Fashion-MNIST.

- [C](#)
- [C++](#)
- [Java](#)
- [Python](#) and [this](#) and [this](#)
- [Scala](#)
- [Go](#)
- [C#](#)
- [NodeJS](#) and [this](#)
- [Swift](#)
- [R](#) and [this](#)
- [Matlab](#)
- [Ruby](#)

Benchmark

We built an automatic benchmarking system based on `scikit-learn` that covers 129 classifiers (but no deep learning) with different parameters. [Find the results here](#).

Name	Parameter	Accuracy (mean)	Accuracy (std)	Training time	Repeats	Job start	Job Done
LinearSVC	{"C":100,"loss":"hinge","multi_class":"crammer_singer","penalty":"l2"}	0.496	0.069	6:50:20	5	3 days ago	2 days ago
LinearSVC	{"C":100,"loss":"hinge","multi_class":"crammer_singer","penalty":"l1"}	0.516	0.059	6:51:28	5	3 days ago	2 days ago
LinearSVC	{"C":100,"loss":"squared_hinge","multi_class":"crammer_singer","penalty":"l1"}	0.492	0.030	7:19:55	5	3 days ago	2 days ago
LinearSVC	{"C":100,"loss":"squared_hinge","multi_class":"crammer_singer","penalty":"l2"}	0.484	0.035	7:17:29	5	3 days ago	2 days ago
GradientBoostingClassifier	{"loss":"deviance","max_depth":50,"n_estimators":10}	0.795	0.001	15:57:02	2	3 days ago	2 days ago
GradientBoostingClassifier	{"loss":"deviance","max_depth":10,"n_estimators":50}	0.872	0.001	12:20:35	2	3 days ago	2 days ago
LinearSVC	{"C":10,"loss":"hinge","multi_class":"crammer_singer","penalty":"l2"}	0.749	0.011	3:55:42	5	3 days ago	2 days ago
LinearSVC	{"C":10,"loss":"hinge","multi_class":"crammer_singer","penalty":"l1"}	0.751	0.028	4:51:51	5	3 days ago	2 days ago
LinearSVC	{"C":10,"loss":"squared_hinge","multi_class":"crammer_singer","penalty":"l2"}	0.748	0.035	4:58:09	5	3 days ago	2 days ago

You can reproduce the results by running `benchmark/runner.py`. We recommend building and deploying [this Dockerfile](#).

You are welcome to submit your benchmark; simply create a new issue and we'll list your results here. Before doing that, please make sure it does not already appear [in this list](#). Visit our [contributor guidelines](#) for additional details.

The table below collects the submitted benchmarks. Note that **we haven't yet tested these results**. You are welcome to validate the results using the code provided by the submitter. Test accuracy may differ due to the number of epoch, batch size, etc. To correct this table, please create a new issue.

Classifier	Preprocessing	Fashion test accuracy	MNIST test accuracy	Submitter	Code
2 Conv+pooling	None	0.876	-	Kashif Rasul	🔗
2 Conv+pooling	None	0.916	-	Tensorflow's doc	🔗
2 Conv+pooling+ELU activation (PyTorch)	None	0.903	-	@AbhirajHinge	🔗
2 Conv	Normalization, random horizontal flip, random vertical flip, random translation, random rotation.	0.919	0.971	Kyriakos Efthymiadis	🔗
2 Conv <100K parameters	None	0.925	0.992	@hardmaru	🔗
2 Conv ~113K parameters	Normalization	0.922	0.993	Abel G.	🔗
2 Conv+3 FC ~1.8M parameters	Normalization	0.932	0.994	@Xfan1025	🔗
2 Conv+3 FC ~500K parameters	Augmentation, batch normalization	0.934	0.994	@cmasch	🔗

Classifier	Preprocessing	Fashion test accuracy	MNIST test accuracy	Submitter	Code
2 Conv+pooling+BN	None	0.934	-	@khanguyen1207	
2 Conv+2 FC	Random Horizontal Flips	0.939	-	@ashmeet13	
3 Conv+2 FC	None	0.907	-	@Cenk Bircanoğlu	
3 Conv+pooling+BN	None	0.903	0.994	@meghanabhang	
3 Conv+pooling+2 FC+dropout	None	0.926	-	@Umberto Griff	
3 Conv+BN+pooling	None	0.921	0.992	@GunjanChhablani	
5 Conv+BN+pooling	None	0.931	-	@Noumanmufc1	
CNN with optional shortcuts, dense-like connectivity	standardization+augmentation+random erasing	0.947	-	@kennivich	
GRU+SVM	None	0.888	0.965	@AFagarap	
GRU+SVM with dropout	None	0.897	0.988	@AFagarap	
WRN40-4 8.9M params	standard preprocessing (mean/std subtraction/division) and augmentation (random crops/vertical flips)	0.967	-	@ajbrock	
DenseNet-BC 768K params	standard preprocessing (mean/std subtraction/division) and augmentation (random crops/vertical flips)	0.954	-	@ajbrock	
MobileNet	augmentation (horizontal flips)	0.950	-	@苏剑林	
ResNet18	Normalization, random horizontal flip, random vertical flip, random translation, random rotation.	0.949	0.979	Kyriakos Efthymiadis	
GoogleNet with cross-entropy loss	None	0.937	-	@Cenk Bircanoğlu	
AlexNet with Triplet loss	None	0.899	-	@Cenk Bircanoğlu	
SqueezeNet with cyclical learning rate 200 epochs	None	0.900	-	@snakers4	
Dual path network with wide resnet 28-10	standard preprocessing (mean/std subtraction/division) and augmentation (random crops/vertical flips)	0.957	-	@Queequeg	
MLP 256-128-100	None	0.8833	-	@heitorrapela	
VGG16 26M parameters	None	0.935	-	@QuantumLiu	

Classifier	Preprocessing	Fashion test accuracy	MNIST test accuracy	Submitter	Code
WRN-28-10	standard preprocessing (mean/std subtraction/division) and augmentation (random crops/horizontal flips)	0.959	-	@zhunzhong07	
WRN-28-10 + Random Erasing	standard preprocessing (mean/std subtraction/division) and augmentation (random crops/horizontal flips)	0.963	-	@zhunzhong07	
Human Performance	Crowd-sourced evaluation of human (with no fashion expertise) performance. 1000 randomly sampled test images, 3 labels per image, majority labelling.	0.835	-	Leo	-
Capsule Network 8M parameters	Normalization and shift at most 2 pixel and horizontal flip	0.936	-	@XifengGuo	
HOG+SVM	HOG	0.926	-	@subalde	
XgBoost	scaling the pixel values to mean=0.0 and var=1.0	0.898	0.958	@anktplwl91	
DENSER	-	0.953	0.997	@fillassuncao	
Dyra-Net	Rescale to unit interval	0.906	-	@Dirk Schäfer	
Google AutoML	24 compute hours (higher quality)	0.939	-	@Sebastian Heinz	

Other Explorations of Fashion-MNIST

[Fashion-MNIST: Year in Review](#)

[Fashion-MNIST on Google Scholar](#)

[Generative adversarial networks \(GANs\)](#)

- Tensorflow implementation of various GANs and VAEs. ([Recommend to read!](#) Note how various GANs generate different results on Fashion-MNIST, which can not be easily observed on the original MNIST.)
- Make a ghost wardrobe using DCGAN
- fashion-mnist的gan玩具
- CGAN output after 5000 steps
- live demo of Generative Adversarial Network model with deeplearn.js
- GAN Playground - Explore Generative Adversarial Nets in your Browser

[Clustering](#)

- Xifeng Guo's implementation of Unsupervised Deep Embedding for Clustering Analysis (DEC)
- Leland McInnes's Uniform Manifold Approximation and Projection (UMAP)

[Video Tutorial](#)

Machine Learning Meets Fashion by Yufeng G @ Google Cloud



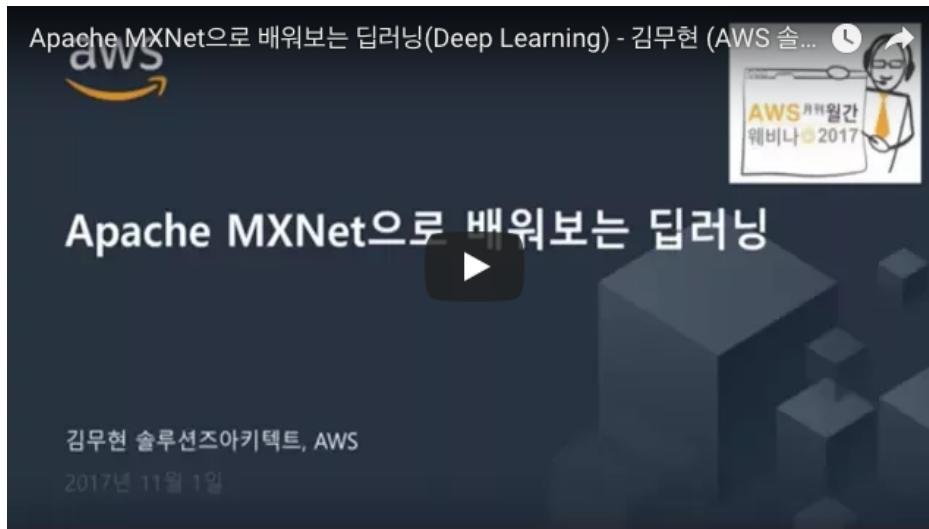
Introduction to Kaggle Kernels by [Yufeng G](#) @ Google Cloud



动手学深度学习 by Mu Li @ Amazon AI



Apache MXNet으로 배워보는 딥러닝(Deep Learning) - 김무현 (AWS 솔루션즈아키텍트)

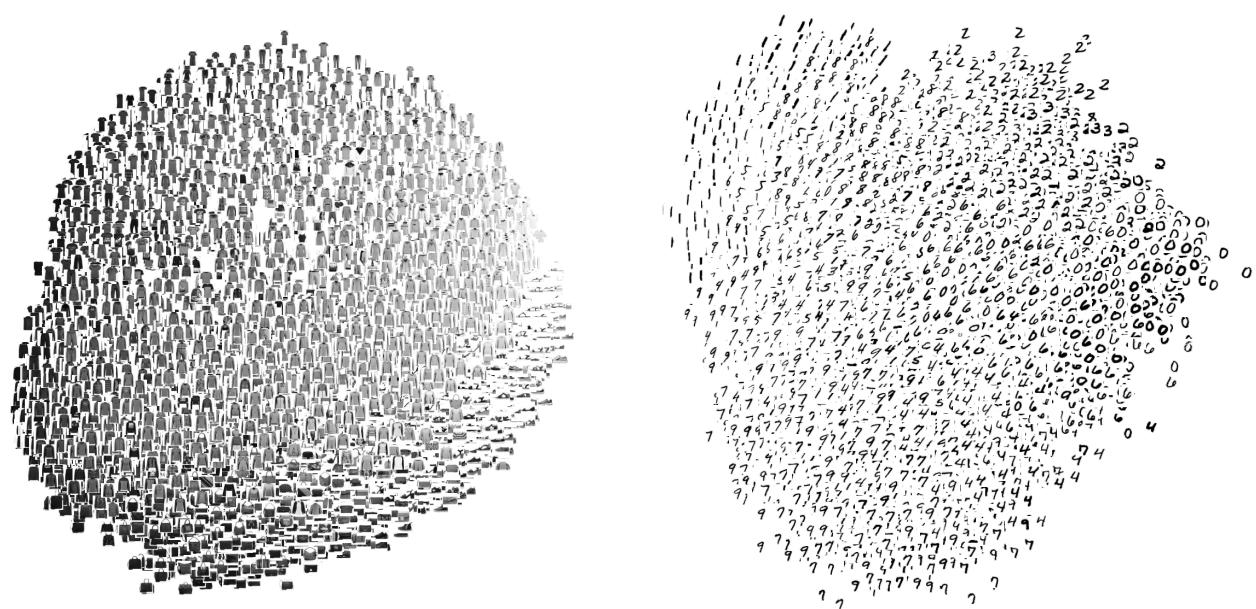


Visualization

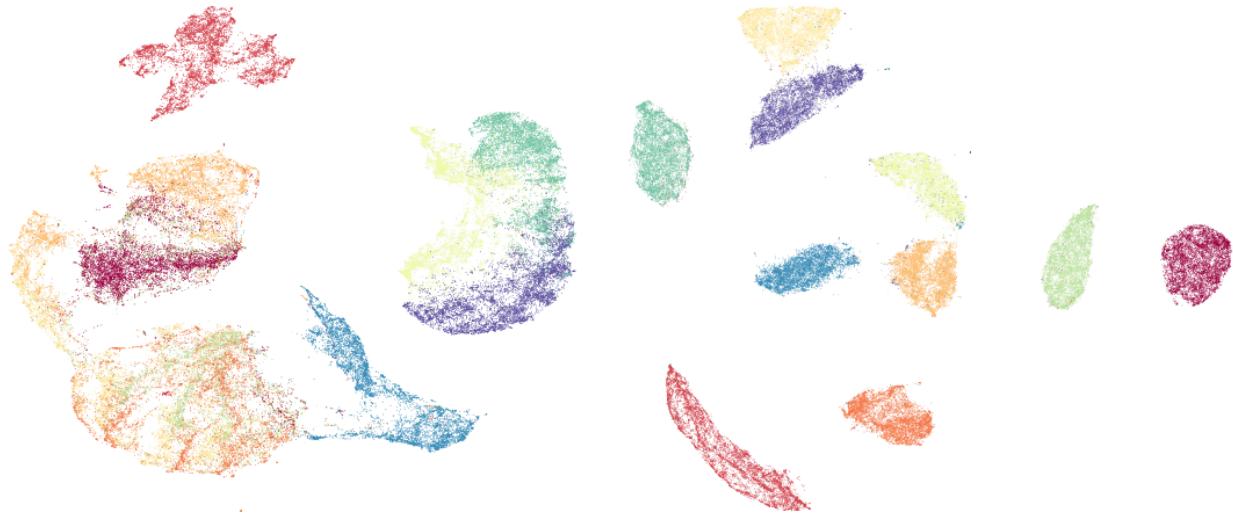
t-SNE on Fashion-MNIST (left) and original MNIST (right)



PCA on Fashion-MNIST (left) and original MNIST (right)



UMAP on Fashion-MNIST (left) and original MNIST (right)



Contributing

Thanks for your interest in contributing! There are many ways to get involved; start with our [contributor guidelines](#) and then check these [open issues](#) for specific tasks.

Contact

To discuss the dataset, please use [chat on gitter](#).

Citing Fashion-MNIST

If you use Fashion-MNIST in a scientific publication, we would appreciate references to the following paper:

Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine Learning Algorithms. Han Xiao, Kashif Rasul, Roland Vollgraf. [arXiv:1708.07747](https://arxiv.org/abs/1708.07747)

Biblatex entry:

```
@online{xiao2017/online,
  author      = {Han Xiao and Kashif Rasul and Roland Vollgraf},
  title       = {Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine Learning Algorithms},
  date        = {2017-08-28},
  year         = {2017},
  eprintclass  = {cs.LG},
  eprinttype   = {arXiv},
  eprint      = {cs.LG/1708.07747},
}
```

[Who is citing Fashion-MNIST?](#)

License

The MIT License (MIT) Copyright © [2017] Zalando SE, <https://tech.zalando.com>

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.